

EMBODIED RETRIEVAL-AUGMENTED CONTROL: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study retrieval-augmented robot policies. The motivating bet is: Retrieve prior physical episodes by mechanism similarity, not language similarity.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For retrieval-augmented robot policies, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026; 2025a;b;c;d;e). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.560	0.440
observed_only_wam	hidden_mode_shift	0.190	0.810
augmentation_as_data_wam	nominal	0.585	0.415
augmentation_as_data_wam	hidden_mode_shift	0.268	0.732
scalar_uncertainty_wam	nominal	0.652	0.348
scalar_uncertainty_wam	hidden_mode_shift	0.320	0.680
proposed_branch_mechanism	nominal	0.780	0.220
proposed_branch_mechanism	hidden_mode_shift	0.613	0.387

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Embodied Retrieval-Augmented Control argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Benchmarking retrieval-augmented large language models in biomedical nlp: Application, robustness, and self-awareness, 2025a. <https://doi.org/10.1126/sciadv.adr1443>.
- Scalable retrieval-augmented generation for technical documents with hallucination risk control, 2025b. <https://doi.org/10.66238/ijcbs29>.
- Advancing large language models with enhanced retrieval-augmented generation: Evidence from biological uav swarm control, 2025c. <https://doi.org/10.3390/drones9050361>.
- Pai-net: Retrieval-augmented generation patent network using prior art information, 2025d. <https://doi.org/10.3390/systems13040259>.
- Enhancing environmental control in broiler production: Retrieval-augmented generation for improved decision-making with large language models, 2025e. <https://doi.org/10.3390/agriengineering7010012>.
- Efficient similarity-based datasheet retrieval and analysis using retrieval-augmented generation for electronic component selection, 2026. <https://doi.org/10.3390/electronics15112301>.